Troubleshooting

Deployment Integration Team

1. Troubleshooting Flow	2
2. Install-config.yaml	5
3. Bootstrap VM issues	6
3.1. Bootstrap VM cannot boot up my cluster nodes	7
3.2. Inspecting Logs	8
4. Cluster nodes won't PXE	10
5. API not accessible	11
6. Review Install	12
7. Following the Installation	
8. Post install pod errors	14
9. NTP out of sync	17
10. Cleaning up previous installations	21
11. Registry Issues	22
12. Miscellaneous Issues	23
12.1. runtime network not ready	23
12.2. Servers not getting the correct IPv6 over DHCP	24
12.3. Servers not getting the correct hostname over DHCP	24

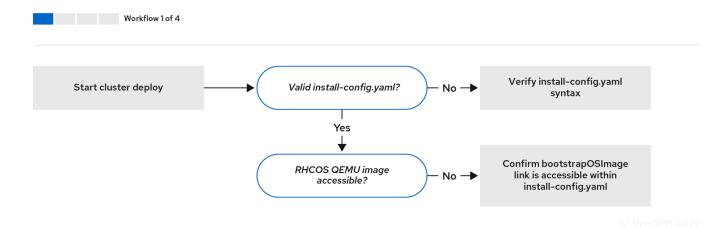


Download the PDF version of this document or visit https://openshift-kni.github.io/baremetal-deploy/

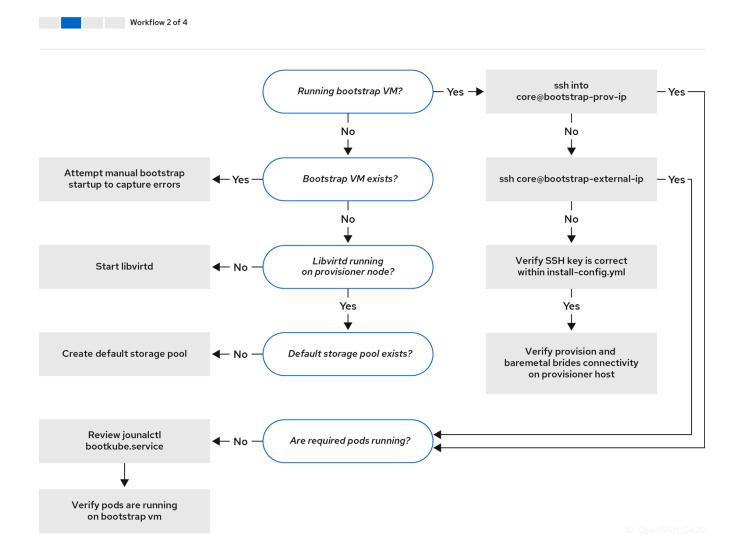
While attempting to deploy Installer Provisioned Infrastructure (IPI) of OpenShift on Bare Metal (BM), you may run into a situation where you need to troubleshoot your environment. This document provides troubleshooting guidance and tips in solving common issues that may arise.

Chapter 1. Troubleshooting Flow

Prior to troubleshooting your environment, it is critical to understand the overall flow of the IPI on BM installation. The diagrams below provide a troubleshooting flow providing a step-by-step breakdown of where to start debugging your environment.

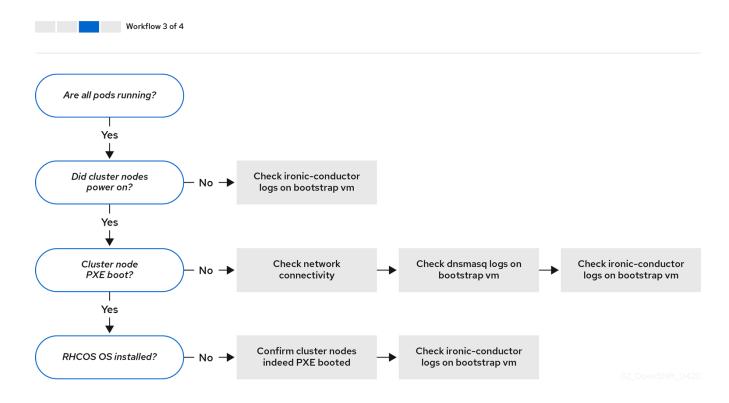


Flow-Diagram-1 relates with issues that relate to your install-config.yaml file. Troubleshooting suggestions can be found at install-config.yaml



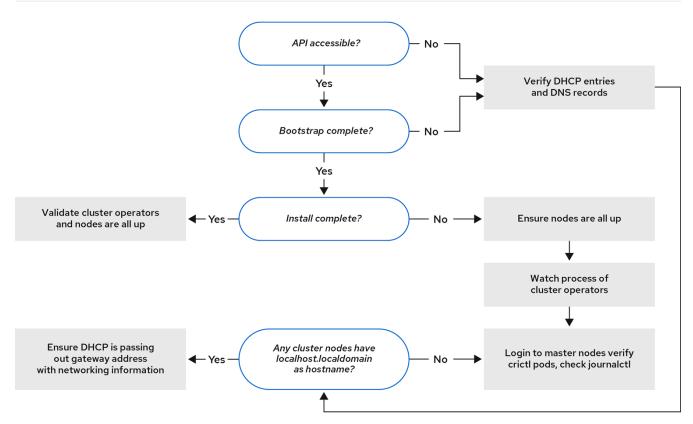
Flow-Diagram-2 relates with issues that relate to your

- Bootstrap VM
- Cluster nodes that won't PXE



Flow-Diagram-3 relates with issues that relate to your

• Cluster nodes that won't PXE



82_OpenShift_0420

Flow-Diagram-4 relates with issues that relate to your

- API accessibility
- reviewing installation

Chapter 2. Install-config.yaml

The install-config.yaml represents all the nodes that will be part of your OpenShift cluster and the necessary options consisting of but not limited to apiVersion, baseDomain, networking, VIPS, imageContentSources etc.

During the deployment of your OpenShift cluster, if errors occur early on due to the install-config.yaml, confirm the following

- Use the guidelines as mentioned in the YAML-tips
- Verifying syntax is correct within your YAML using syntax-check
- Verify the RHCOS QEMU images are properly defined and accessible via the URL provided in the install-config.yaml via a command such as:

```
curl -s -o /dev/null -I -w "%{http_code}\n"
http://webserver.example.com:8080/rhcos-44.81.202004250133-0-
qemu.x86_64.qcow2.gz?sha256=7d884b46ee54fe87bbc3893bf2aa99af3b2d31f2e19ab5529c60636
fbd0f1ce7
```

The output of 200 ensures that we are geting a valid response from our webserver that is indeed storing our bootstrap VM image in the above example.

Chapter 3. Bootstrap VM issues

The bootstrap node is a virtual machine that is triggered by the OpenShift installer that handles the provisioning of the nodes that will be part of the OpenShift cluster.

Once the deployment is triggered (may take 10 to 15 minutes) we can see if the Bootstrap VM is operational. We can verify the bootstrap VM is operational via the virsh command.



The name of the bootstrap VM is always the clustername, followed by random set of characters and ending in the word bootstrap.

If no bootstrap VM is running after waiting 10-15 minutes, it is important to troubleshoot as to why it was not created. Possible issues to verify are:

• Verify libvirtd is running on the system

If the bootstrap VM is indeed operational, we want to verify if we can log into it.

We can find the IP address of the bootstrap VM using the virsh console command as shown.

```
[kni@provisioner ~]$ sudo virsh console kni7-4q49b-bootstrap
Connected to domain kni7-4q49b-bootstrap
Escape character is ^]

Red Hat Enterprise Linux CoreOS 43.81.202001142154.0 (Ootpa) 4.3
SSH host key: SHA256:BRWJktXZgQQRY5zjuAV0IKZ4WM7i4TiUyMVanqu9Pqg (ED25519)
SSH host key: SHA256:7+iKGA7VtG5szmk2jB5gl/5EZ+SNcJ3a2g23o0lnIio (ECDSA)
SSH host key: SHA256:DH5VWhvhvagOTaLsYiVNse9ca+ZSW/3000Med8rIGOc (RSA)
ens3: fd35:919d:4042:2:c7ed:9a9f:a9ec:7
ens4: 172.22.0.2 fe80::1d05:e52e:be5d:263f
localhost login:
```

• Once the IP address has been obtained, we can attempt to login to the bootstrap VM using the ssh command.



In the console output above, the IPv6 IP provided by ens3 or the IPv4 IP provided by ens4 can be used.

```
[kni@provisioner ~]$ ssh core@172.22.0.2
```

If attempting to log into the bootstrap VM and you are not successful, you've likely reached one of the following scenarios:

- Cannot reach the 172.22.0.0/24 network. This means you should verify your connectivity on the provisioning host specifically around the provisioning network bridge
- Cannot reach the bootstrap VM via the public network. When attempting to SSH via baremetal network you should verify your connectivity on the provisioning host specifically around the baremetal network bridge.
- Permission denied (publickey,password,keyboard-interactive). When attempting to access the bootstrap VM an error of Permission denied might occur, verify that your SSH key for the user attempting to log into the VM is set within your install-config.yaml file.

3.1. Bootstrap VM cannot boot up my cluster nodes

During the deployment, it is possible for the bootstrap VM to not boot up the OpenShift cluster nodes in order to provision them with their respective RHCOS OS image. This issue can arise due to:

- Problem with the install-config.yaml file
- · Issues with Out of band network access from the baremetal network

To verify the issue, there are 3 basic containers related to ironic

- * Ironic-api
- * Ironic-conductor
- * Ironic-inspector

The ironic-conductor pod includes the most details regarding the attempted boot up of your OpenShift cluster nodes as it attempts to perform the login and execute the action over IPMI, you can check the logs as per the below syntax

```
# podman logs -f ironic-conductor
```

If you encounter an issue where the master nodes are not booting up via PXE, check the ironic-conductor pod

Potential reason

They might be in ON state when deployment started.

Solution

Make sure to power off the OpenShift nodes before you begin the installation via iDRAC/iLO.

```
ipmitool -I lanplus -U root -P <Password> -H <Out-of-band-ip> power off
```

3.2. Inspecting Logs

• bootstrap VM cannot download RHCOS image

When experiencing issues downloading or access the RHCOS images, verify the following two containers are up and running:

```
ipa-downloader
coreos-downloader
```

These containers download resources from a provided webserver or the external quay.io registry (which ever is provided in the install-config.yaml).

Example of internal webserver hosting RHCOS images

```
bootstrapOSImage: http://<ip:port>/rhcos-43.81.202001142154.0-
qemu.x86_64.qcow2.gz?sha256=9d999f55ff1d44f7ed7c106508e5deecd04dc3c06095d34d36bf1cd127
837e0c
clusterOSImage: http://<ip:port>/rhcos-43.81.202001142154.0-
openstack.x86_64.qcow2.gz?sha256=a1bda656fa0892f7b936fdc6b6a6086bddaed5dafacedcd7a1e81
1abb78fe3b0
```

Check the status of these containers using

```
[core@localhost ~]$ podman logs -f ipa-downloader
[core@localhost ~]$ podman logs -f coreos-downloader
```

- There can be scenarios where there is a problem in the baremetal network that uses external registry (https://quay.io) to download the images, make use of curl command to confirm that you can reach the registry.
- If using a caching webserver to host the images, bootstrap VM cannot access the URL and verify it is correct.

With the ability to login to the bootstrap VM

```
[kni@provisioner ~]$ ssh core@172.22.0.2
```

Issue the following journalctl commands to inspect the bootkube logs that shows all the containers launched during the deployment phase.

Within the bootstrap VM,

```
[core@localhost ~]$ journalctl -xe
[core@localhost ~]$ journalctl -b -f -u bootkube.service
```

Verify all the pods(dnsmasq, mariadb, httpd, ironic pods etc) are in running state using sudo podman ps command (displays the containers).

If there are issues with the pods, check the logs of all the containers with issues. Use the below command for an example on how to check the log of the ironic-apis:

Example one

```
[core@localhost ~]$ sudo podman logs <container-id>
```

Chapter 4. Cluster nodes won't PXE

When dealing with issues of OpenShift cluster nodes not PXE booting the following issues should be further verified:

- Check the network connectivity to your Provisioning Network
- Make sure to have PXE enabled on the provisioning NIC and PXE disabled for all other NICs.
- Verify you have the proper boot MAC address(provisioning NIC MAC address) and hardware profiles in install-config.yaml file.

Master node settings

bootMACAddress: 24:6E:96:1B:96:90 # MAC of bootable provisioning NIC

hardwareProfile: default #master node settings

Worker node settings

bootMACAddress: 24:6E:96:1B:96:90 # MAC of bootable provisioning NIC

hardwareProfile: unknown #worker node settings

Chapter 5. API not accessible

When the API is not accessible these issues often arise when the following is not properly addressed.

- Incorrect name resolutions are set in the DNS server, make use of dig and nslookup commands.
- SSH into the OpenShift cluster and make sure they are getting the FQDN and not just localhost.localdomain.

For example, if we wanted to verify the API VIP was properly set via dig, we could:

```
dig api.<cluster-name>.example.com
; <<>> DiG 9.11.4-P2-RedHat-9.11.4-26.P2.el8 <<>> api.<cluster-name>.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 37551
;; flags: qr aa rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 2
;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
; COOKIE: 866929d2f8e8563582af23f05ec44203d313e50948d43f60 (good)
;; QUESTION SECTION:
;api.<cluster-name>.example.com. IN A
;; ANSWER SECTION:
api.<cluster-name>.example.com. 10800 IN A 10.19.13.86
;; AUTHORITY SECTION:
<cluster-name>.example.com. 10800 IN NS <cluster-name>.example.com.
;; ADDITIONAL SECTION:
<cluster-name>.example.com. 10800 IN A 10.19.14.247
;; Query time: 0 msec
;; SERVER: 10.19.14.247#53(10.19.14.247)
;; WHEN: Tue May 19 20:30:59 UTC 2020
;; MSG SIZE rcvd: 140
```

This concludes that the appropriate IP for the api.<cluster-name>.example.com VIP was 10.19.13.86. This IP address should reside on the baremetal network.

Chapter 6. Review Install

When the OpenShift cluster nodes are installed appropriately, the following Ready state is seen within the STATUS column.

```
[kni@provisioner ~]$ oc get nodes
                      STATUS
                               ROLES
                                              AGE VERSION
master-0.example.com
                      Ready
                               master,worker
                                              4h
                                                   v1.16.2
master-1.example.com
                      Ready
                                              4h
                                                   v1.16.2
                               master,worker
master-2.example.com
                      Ready
                               master,worker
                                              4h
                                                   v1.16.2
```

Confirm all pods are successfully deployed. The following command removes any pods that are still running or have completed as part of the output.

```
[kni@provisioner ~]$ oc get pods --all-namespaces | grep -iv running | grep -iv complete
```

Chapter 7. Following the Installation

During the deployment process, we can check its overall status by issuing the tail command to the log file .openshift_install.log which will be under the install directory folder.

tail -f /path/to/install-dir/.openshift_install.log

Chapter 8. Post install pod errors

The following section describes how to handle metal3 pod errors during deployment. (only applies for 4.3)

```
[kni@provisioner ~]$ oc get all -n openshift-machine-api
NAME READY STATUS
RESTARTS AGE
pod/metal3-6c6cc7c56c-lj4lr 0/8 Init:CreateContainerConfigError
0 <invalid>
```

This happens due to the missing ConfigMap

```
Warning Failed <invalid> (x7 over 32s) kubelet, master-1.<cluster-name>.example.com Error: configmap "metal3-config" not found
```

Make sure to add the ConfigMap before starting the deployment as shown below to solve the above problem

Create your metal3-config.yaml file as shown below:

```
apiVersion: v1
kind: ConfigMap
metadata:
   name: metal3-config
   namespace: openshift-machine-api
data:
   cache_url: rhcos-43.81.202001142154.0-qemu.x86_64.qcow2.gz
   deploy_kernel_url: http://172.22.0.1:6180/images/ironic-python-agent.kernel
   deploy_ramdisk_url: http://172.22.0.1:6180/images/ironic-python-agent.initramfs
   dhcp_range: 172.22.0.10,172.22.0.100
   http_port: "6180"
   ironic_endpoint: http://172.22.0.1:6385/v1/
   ironic_inspector_endpoint: http://172.22.0.3:5050/v1/
   provisioning_interface: eno1
   provisioning_ip: 172.22.0.1/24
   rhcos_image_url: <URL-which-has-qcow-image>
```

Note: Change your rhcos_image_url to the appropriate URL for your deployment.

Place this file in <Install-Dir>

```
[kni@provisioner ~]$ cp metal3-config.yaml ocp/openshift/99_metal3-config.yaml
```

and then re-run the installation to fix the metal³ pod issue.

After the installation is complete copy the config file to the .kube/config directory to interact with the cluster.

```
[kni@provisioner ~]$ export KUBECONFIG=/install-dir/ocp/auth/kubeconfig
```

Verify if all the OpenShift nodes are up and running

```
[kni@provisioner ~]$ oc get nodes
                                         STATUS
NAME
                                                  ROLES
                                                                   AGE
                                                                          VERSION
master-0.cloud.lab.eng.bos.redhat.com
                                         Ready
                                                  master, worker
                                                                   4h
                                                                        v1.16.2
master-1.cloud.lab.eng.bos.redhat.com
                                                                        v1.16.2
                                         Ready
                                                  master, worker
                                                                   4h
master-2.cloud.lab.eng.bos.redhat.com
                                         Ready
                                                  master, worker
                                                                        v1.16.2
                                                                   4h
```

If the installation fails:

- Please check for logs .openshift-install.log.
- Master Node in NotReady state.

If any of the OpenShift nodes are in NotReady state, make sure the kubelet services are enabled.

```
[kni@provisioner~]$ oc get nodes
                             STATUS
                                         ROLES
                                                  AGE
                                                        VERSION
master-0.cloud.example.com
                             NotReady
                                                  67m
                                                        v1.16.2
                                         master
master-1.cloud.example.com
                             Ready
                                         master
                                                  67m
                                                        v1.16.2
master-2.cloud.example.com
                             Ready
                                         master
                                                  67m
                                                        v1.16.2
```

Verify kubelet on each master node, example of logging in to master-0 below

Start the kubelet services and check back the status of the node.

```
[core@master-0 ~]$ sudo systemctl start kubelet
[core@master-0 ~]$ sudo systemctl status kubelet

    kubelet.service - Kubernetes Kubelet

   Loaded: loaded (/etc/systemd/system/kubelet.service; enabled; vendor preset:
enabled)
  Drop-In: /etc/systemd/system/kubelet.service.d
           ──10-default-env.conf, 20-nodenet.conf
   Active: active (running) since Tue 2020-03-10 16:07:27 UTC; 4s ago
[kni@provisioner~]$ oc get nodes
                             STATUS
                                      ROLES
                                                     VERSION
                                               AGE
master-0.cloud.example.com
                             Ready
                                               80m
                                                     v1.16.2
                                      master
master-1.cloud.example.com
                             Ready
                                               80m
                                                     v1.16.2
                                      master
master-2.cloud.example.com
                             Ready
                                               80m
                                                     v1.16.2
                                      master
```

Chapter 9. NTP out of sync

The deployment of OpenShift clusters depends on having synced time. Without synced time, the deployment may fail due to the time inaccuracies if larger than two seconds. For example, notice the AGE of the nodes below.

```
[kni@provisioner ~]$ oc get nodes
NAME
                             STATUS
                                      ROLES
                                               AGE
                                                      VERSION
master-0.cloud.example.com
                                                       v1.16.2
                             Ready
                                      master
                                                145m
master-1.cloud.example.com
                             Ready
                                      master
                                                135m
                                                       v1.16.2
master-2.cloud.example.com
                             Ready
                                      master
                                                145m
                                                      v1.16.2
worker-2.cloud.example.com
                             Ready
                                      worker
                                                100m v1.16.2
```

This leads to inconsistent timing delays as shown below

In an already existing cluster, this can be fixed by following the below steps.

• Add chrony.conf file and encode it as base64, for example:

```
cat << EOF | base 64
server <NTP-server> iburst
stratumweight 0
driftfile /var/lib/chrony/drift
rtcsync
makestep 10 3
bindcmdaddress 127.0.0.1
bindcmdaddress ::1
keyfile /etc/chrony.keys
commandkey 1
generatecommandkey
noclientlog
logchange 0.5
logdir /var/log/chrony
EOF
[text-in-base-64]
```

• Create the MachineConfig file, replacing the base64 string with the one you just created yourself. This example adds the file to master nodes. You can change it to worker or make an additional MachineConfig for the worker role:

```
[kni@provisioner ~]$ cat << EOF > ./99_masters-chrony-configuration.yaml
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  creationTimestamp: null
  labels:
    machineconfiguration.openshift.io/role: master
  name: 99-master-etc-chrony-conf
spec:
  config:
    ignition:
      config: {}
      security:
        tls: {}
      timeouts: {}
      version: 2.2.0
    networkd: {}
    passwd: {}
    storage:
      files:
      - contents:
          source: data:text/plain;charset=utf-8;base64,[text-in-base-64]
          verification: {}
        filesystem: root
        group:
          name: root
        mode: 420
        path: /etc/chrony.conf
        user:
          name: root
    systemd: {}
  osImageURL: ""
```

Make a backup copy of the configuration file.

Since, the cluster is already running, apply the file as follows:



When wanting to setup time sync prior to deployment, generate manifest files, add this file to the openshift directory, then continue to create the cluster.

[kni@provisioner \sim]\$ cp chrony-masters.yaml <Install-Dir>/ocp/openshift/99_masters-chrony-configuration.yaml

Chapter 10. Cleaning up previous installations

If re-running the deployment, delete all the older bootstrap VMs including its volume as shown below. Failure to do so may cause over consumption of resources on the provisioner host.

Clean up the older VMs and their volumes with the help of below script.

```
BOOTSTRAP=$(virsh list --all | grep "$CLUSTER.*bootstrap" | tail -1 | awk '{print $2}')
virsh destroy $BOOTSTRAP
virsh undefine $BOOTSTRAP
virsh vol-delete $BOOTSTRAP default
virsh vol-delete $BOOTSTRAP.ign default
```

Chapter 11. Registry Issues

At times, there may be a need to create a disconnected registry due to limited connectivity to an environment or creating a central location to host images locally. Details can be found in Creating a disconnected registry

When attempting to mirror the registry and if you receive a 'User Not Authorized' error, this is due to not appending the new authentication to your already existing pull-secret.txt file.

After the registry is mirrored, confirm that you can access this in your disconnected environment

```
[kni@provisioner ~]$ curl -k -u <user>:<password>
https://registry.example.com:<registry-port>/v2/_catalog
{"repositories":["<Repo-Name>"]}
```

Chapter 12. Miscellaneous Issues

12.1. runtime network not ready

After the deployment of a cluster if you receive such an error runtime network not ready: NetworkReady=false reason:NetworkPluginNotReady message:Network plugin returns error: Missing CNI default network we need to inspect the pods in openshift-network-operator namespace.

```
[kni@provisioner ~]$ oc get all -n openshift-network-operator

NAME READY STATUS RESTARTS AGE

pod/network-operator-69dfd7b577-bg89v 0/1 ContainerCreating 0 149m
```

The cluster network operator is responsible for deploying the networking components. It does this in response to a special object created by the installer.

It runs very early in the installation process, after the master nodes have come up but before the bootstrap control plane has been torn down. It can be indicative of more subtle installer issues, such as long delays in bringing up master nodes or issues with apiserver communication.

First, determine that the network configuration exists, from provisioner node:

```
[kni@provisioner ~]$ kubectl get network.config.openshift.io cluster -oyaml
apiVersion: config.openshift.io/v1
kind: Network
metadata:
   name: cluster
spec:
   serviceNetwork:
   - 172.30.0.0/16
   clusterNetwork:
   - cidr: 10.128.0.0/14
     hostPrefix: 23
   networkType: OpenShiftSDN
```

If it doesn't exist, the installer didn't create it. You'll have to run openshift-install create manifests to determine why.

Next, check that the network-operator is running, from provisioner node issue:

```
[kni@provisioner ~]$ kubectl -n openshift-network-operator get pods
```

and retrieve the logs. Note that, on multi-master systems, the operator will perform leader election and all other operators will sleep:

```
[kni@provisioner ~]$ kubectl -n openshift-network-operator logs -l <mark>"name=network-operator"</mark>
```

Note: For more details please refer Troubleshooting

12.2. Servers not getting the correct IPv6 over DHCP

- 1. The reserved IPv6 addresses should be outside your DHCP Range
- 2. In the reservation make sure you're using the correct MAC Address and Client ID (if supported), example below:

```
# This is a dnsmasq dhcp reservation, 'id:00:03:00:01' is the client id and '18:db:f2:8c:d5:9f' is the MAC Address for the NIC id:00:03:00:01:18:db:f2:8c:d5:9f,openshift-master-1,[2620:52:0:1302::6]
```

3. Make sure Route Announcements are working and that DHCP server is listening on the required interfaces serving the required ranges

12.3. Servers not getting the correct hostname over DHCP

During the IPv6 deployment you need your servers to get their hostname over DHCP and at times the hostname is not assigned by NetworkManager right away.

If you login in the master nodes and you find something like this:

```
Failed Units: 2
NetworkManager-wait-online.service
nodeip-configuration.service
```

The node likely booted up without hostname, which cause kubelet to boot with localhost.localdomain hostname. The following steps can be taken to attempt to remedy.

1. Force the hostname renewal

```
# Check hostname, if it's localhost run below steps.
hostname
# Wired Connection 5 corresponds to the connection which is assigned to the NIC
connected to the baremetal network, it may be different in your env
# This will force the host to renew the dhcp lease
[core@master-X ~]$ sudo nmcli con up "Wired connection 5"
# Check hostname again
[core@master-X ~]$ hostname
# If the hostname is still localhost.localdomain restart NetworkManager
[core@master-X ~]$ sudo systemctl restart NetworkManager
# If the hostname is still localhost.localdomain wait a few seconds/minutes and
check again, if same, repeat previous steps
```

2. Restart nodeip-configuration service

```
# This service will reconfigure Kubelet service with the correct hostname references
[core@master-X ~]$ sudo systemctl restart nodeip-configuration.service
```

3. Restart kubelet service

```
# We need to reload the unit files definition since kubelet one was changed by previous step sudo systemctl daemon-reload # Restart kubelet [core@master-X ~]$ sudo systemctl restart kubelet.service
```

4. Ensure kubelet booted with correct hostname

```
# Tail the journal and find the hostname
[core@master-X ~]$ sudo journalctl -fu kubelet.service
```



where X is the master node number.

If you already had this cluster running when the node booted up with the wrong hostname you will have pending csrs on the cluster, you **MUST NOT** approve them, therewise you will face even more issues.

```
# Get CSR on the cluster
[kni@provisioner ~]$ oc get csr
# If you see Pending csr ensure they are good ones
[kni@provisioner ~]$ oc get csr <pending_csr> -o jsonpath='{.spec.request}' | base64
-d | openssl req -noout -text
# If you see Subject Name: localhost.localdomain remove the csr
[kni@provisioner ~]$ oc delete csr <wrong_csr>
```